



# CHAPTER 6

## Basic SQL

# Chapter 6 Outline

- SQL Data Definition and Data Types
- Specifying Constraints in SQL
- Basic Retrieval Queries in SQL
- INSERT, DELETE, and UPDATE Statements in SQL
- Additional Features of SQL

# Basic SQL

## ■ SQL language

- Considered one of the major reasons for the commercial success of relational databases
- SQL1 (SQL-86), SQL2 (SQL-92), SQL3 (SQL-99), SQL2003, SQL2006, SQL2008, SQL2011

## ■ SQL

- The origin of SQL is **relational predicate calculus** called tuple calculus (see Ch.8) which was proposed initially as the language **SQUARE** (“**Specifying queries as relational expressions**”).
- **SQL** Actually comes from the word “**SEQUEL**” (**Structured English QUery Language**) which was the original term used in the paper: “SEQUEL TO SQUARE” by Chamberlin and Boyce. IBM could not copyright that term, so they abbreviated to SQL and **copyrighted the term SQL**.
- Now popularly known as “Structured Query language”.
- SQL is an **informal or practical rendering** of the relational data model with syntax

# SQL Standards

- SQL has gone through many standards: starting with **SQL-86** or **SQL 1**. A revised and expanded standard, **SQL-92**, is referred to as **SQL-2**.
- Later standards (from **SQL-1999**) are divided into **core specification** and **specialized extensions**. The extensions are implemented for different applications – such as data mining, data warehousing, multimedia etc.
- SQL-2006 added XML features (Ch. 13); In 2008 they added Object-oriented features (Ch. 12). Further update is SQL-2011.
- **SQL-3** is the current standard which started with **SQL-1999**. It is not fully implemented in any RDBMS.

# SQL Data Definition, Data Types, Standards

- Terminology:
  - **Table**, **row**, and **column** used for relational model terms relation, tuple, and attribute
- CREATE statement
  - Main SQL command for data definition
- The language has features for : **Data definition**, **Data Manipulation**, Transaction control (Transact-SQL, Ch. 20), **Indexing** (Ch.17), Security specification (Grant and Revoke- see Ch.30), Active databases (Ch.26), Multi-media (Ch.26), Distributed databases (Ch.23) etc.
  - SQL is both a DDL and a DML

# Schema and Catalog Concepts in SQL

- We cover the **basic standard SQL** syntax – there are variations in existing RDBMS systems
- **SQL schema**
  - Identified by a **schema name**
  - Includes an **authorization identifier** and **descriptors** for each element
- **Schema elements** include
  - Tables, constraints, views, domains, and other constructs
- Each statement in SQL ends with a **semicolon**

# Schema and Catalog Concepts in SQL (cont'd.)

- **CREATE SCHEMA statement**
  - **CREATE SCHEMA** COMPANY **AUTHORIZATION** 'Jsmith' ;
- **Catalog**
  - Named collection of schemas in an SQL environment
- SQL also has the concept of a cluster of catalogs.
- **SQL environment**
  - Installation of an **SQL-compliant** RDBMS on a computer system



# The **CREATE TABLE** Command in SQL

- Specifying a new relation
  - Provide name of table
  - Specify attributes, their types and initial constraints
- Can optionally specify schema:
  - **CREATE TABLE** COMPANY.EMPLOYEE . . .  
or
  - **CREATE TABLE** EMPLOYEE . . .

# The CREATE TABLE Command in SQL (cont'd.)

- **Base tables (base relations)**
  - Relation and its tuples are actually created and stored as a file by the DBMS
- **Virtual relations (views)**
  - Created through the **CREATE VIEW** statement. Do not correspond to any physical file.

# COMPANY relational database schema (Fig. 5.7)

## EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

## DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------

## DEPT\_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
----------------	------------------

## PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

## WORKS\_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

## DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------

# One possible database state for the COMPANY relational database schema (Fig. 5.6)

## EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

## DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

## DEPT\_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston

# One possible database state for the COMPANY relational database schema – continued (Fig. 5.6)

**WORKS\_ON**

<u>Essn</u>	<u>Pno</u>	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

**PROJECT**

<u>Pname</u>	<u>Pnumber</u>	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

**DEPENDENT**

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
333445555	Alice	F	1986-04-05	Daughter
333445555	Theodore	M	1983-10-25	Son
333445555	Joy	F	1958-05-03	Spouse
987654321	Abner	M	1942-02-28	Spouse
123456789	Michael	M	1988-01-04	Son
123456789	Alice	F	1988-12-30	Daughter
123456789	Elizabeth	F	1967-05-05	Spouse

# SQL CREATE TABLE data definition statements for defining the COMPANY schema from Figure 5.7 (Fig. 6.1)

## **CREATE TABLE EMPLOYEE**

( Fname	VARCHAR(15)	NOT NULL,
Minit	CHAR,	
Lname	VARCHAR(15)	NOT NULL,
Ssn	CHAR(9)	NOT NULL,
Bdate	DATE,	
Address	VARCHAR(30),	
Sex	CHAR,	
Salary	DECIMAL(10,2),	
Super_ssn	CHAR(9),	
Dno	INT	NOT NULL,

**PRIMARY KEY** (Ssn),

## **CREATE TABLE DEPARTMENT**

( Dname	VARCHAR(15)	NOT NULL,
Dnumber	INT	NOT NULL,
Mgr_ssn	CHAR(9)	NOT NULL,
Mgr_start_date	DATE,	

**PRIMARY KEY** (Dnumber),

**UNIQUE** (Dname),

**FOREIGN KEY** (Mgr\_ssn) **REFERENCES** EMPLOYEE(Ssn) );

## **CREATE TABLE DEPT\_LOCATIONS**

( Dnumber	INT	NOT NULL,
Dlocation	VARCHAR(15)	NOT NULL,

**PRIMARY KEY** (Dnumber, Dlocation),

**FOREIGN KEY** (Dnumber) **REFERENCES** DEPARTMENT(Dnumber) );

*continued on next slide*

# SQL CREATE TABLE data definition statements for defining the COMPANY schema from Figure 5.7 (Fig. 6.1)-continued

**CREATE TABLE PROJECT**

( Pname	VARCHAR(15)	NOT NULL,
Pnumber	INT	NOT NULL,
Plocation	VARCHAR(15),	
Dnum	INT	NOT NULL,

**PRIMARY KEY** (Pnumber),

**UNIQUE** (Pname),

**FOREIGN KEY** (Dnum) **REFERENCES** DEPARTMENT(Dnumber) );

**CREATE TABLE WORKS\_ON**

( Essn	CHAR(9)	NOT NULL,
Pno	INT	NOT NULL,
Hours	DECIMAL(3,1)	NOT NULL,

**PRIMARY KEY** (Essn, Pno),

**FOREIGN KEY** (Essn) **REFERENCES** EMPLOYEE(Ssn),

**FOREIGN KEY** (Pno) **REFERENCES** PROJECT(Pnumber) );

**CREATE TABLE DEPENDENT**

( Essn	CHAR(9)	NOT NULL,
Dependent_name	VARCHAR(15)	NOT NULL,
Sex	CHAR,	
Bdate	DATE,	
Relationship	VARCHAR(8),	

**PRIMARY KEY** (Essn, Dependent\_name),

**FOREIGN KEY** (Essn) **REFERENCES** EMPLOYEE(Ssn) );

# The CREATE TABLE Command in SQL (cont'd.)

- Some foreign keys may cause errors
  - Specified either via:
    - Circular references
      - Ex: EMPLOYEE.Super\_ssn refers EMPLOYEE itself
    - Or because they refer to a table that has not yet been created
      - EMPLOYEE.Dno refers to DEPARTMENT
- DBA's have ways to stop referential integrity enforcement to get around this problem.
  - The constraints are left out of the initial CREATE TABLE statement and then added later using the ALTER TABLE statement



# Attribute Data Types and Domains in SQL

- Basic data types

- **Numeric data types**

- Integer numbers: INTEGER, INT, and SMALLINT
    - Floating-point (real) numbers: FLOAT or REAL, and DOUBLE PRECISION
    - Format is DECIMAL(*i*, *j*) or DEC(*i*, *j*) or NUMERIC(*i*, *j*) where *i* is the total number of decimal digits (the **precision**) and *j* is the number of digits after the decimal point (the **scale**)
      - The default for **scale** is zero, and the default for **precision** is implementation-defined

# Attribute Data Types and Domains in SQL (cont'd.)

## ■ Character-string data types

- Fixed length: `CHAR (n)` , `CHARACTER (n)`
  - A shorter string is padded with blank characters to the right
- Varying length: **`VARCHAR (n)`** , `CHAR VARYING (n)` , `CHARACTER VARYING (n)` where *n* is the maximum number of characters
- The literal string value is placed between single quotation and is **case sensitive**. E.g., 'yourString'
- **Concatenation operator denoted by ||** that can concatenate two string
  - E.g., 'abc'||'XYZ' results in a single string 'abcXYZ'
- Large text values, such as document: **`CHARACTER LARGE OBJECT`** or **`CLOB`**
  - **`CLOB (20M)`**

# Attribute Data Types and Domains in SQL (cont'd.)

- **Bit-string data types**

- Fixed length: `BIT (n)`

- The literal bit strings are placed between single but preceded by a **B** to distinguish them from character string. E.g., **B'10101'**

- Large binary values, such as image: **BINARY LARGE OBJECT (or BLOB)**

- **BLOB(30G)**

- Varying length: `BIT VARYING (n)`

- **Boolean data type**

- Values of `TRUE` or `FALSE` or `NULL`

- **DATE data type**

- **Ten positions**

- Components are `YEAR`, `MONTH`, and `DAY` in the form **YYYY-MM-DD**

- Multiple mapping functions available in RDBMSs to change date formats

# Attribute Data Types and Domains in SQL (cont'd.)

- **TIME:**

- Made up of **hour:minute:second** in the format **hh:mm:ss**

- **TIME(i):** time fractional seconds precision

- Made up of **hour:minute:second** plus **i** additional digits specifying fractions of a second
- format is **hh:mm:ss:ii...i**

- **TIME WITH TIME ZONE:**

- includes an additional six positions for specifying the displacement from the standard universal time zone, which is in the range **+13:00 ---- -12:59**

# Attribute Data Types and Domains in SQL (cont'd.)

- Additional data types
  - **Timestamp** data type
    - Includes the `DATE` and `TIME` fields
    - Plus a minimum of six positions for decimal fractions of seconds
    - E.g., `TIMESTAMP '2008-09-27 09:12:47.648302'`
    - Optional `WITH TIME ZONE` qualifier
  - **INTERVAL** data type
    - Specifies a relative value that can be used to increment or decrement an absolute value of a date, time, or timestamp
    - Intervals are qualified to be either `YEAR-MONTH` intervals or `DAY-TIME` intervals
  - **DATE**, **TIME**, **Timestamp**, **INTERVAL** data types can be **cast** or converted to string formats for comparison.

# INTERVAL Data Type

Type identifier	Class	Description
MONTH	Year-Month	Number of months between two dates.
YEAR	Year-Month	Number of years between two dates.
YEAR_TO_MONTH	Year-Month	Number of years and months between two dates.
DAY	Day-Time	Number of days between two dates.
HOUR	Day-Time	Number of hours between two date/times.
MINUTE	Day-Time	Number of minutes between two date/times.
SECOND	Day-Time	Number of seconds between two date/times.
DAY_TO_HOUR	Day-Time	Number of days/hours between two date/times.
DAY_TO_MINUTE	Day-Time	Number of days/hours/minutes between two date/times.
DAY_TO_SECOND	Day-Time	Number of days/hours/minutes/seconds between two date/times.
HOUR_TO_MINUTE	Day-Time	Number of hours/minutes between two date/times.
HOUR_TO_SECOND	Day-Time	Number of hours/minutes/seconds between two date/times.
MINUTE_TO_SECOND	Day-Time	Number of minutes/seconds between two date/times.

```
CREATE TABLE test (  
  id    DECIMAL PRIMARY KEY,  
  col1  INTERVAL YEAR TO MONTH,  
  col2  INTERVAL DAY TO SECOND(6)  
);
```

# Attribute Data Types and Domains in SQL (cont'd.)

## ■ Domain

- Name used with the attribute specification
- Makes it easier to change the data type for a domain that is used by numerous attributes
- Improves schema readability
- Example:
  - **CREATE DOMAIN** SSN\_TYPE **AS** CHAR(9) ;

## ■ TYPE

- User Defined Types (UDTs) are supported for object-oriented applications. (See Ch.12) Uses the command: **CREATE TYPE**

# Specifying Constraints in SQL

## Basic constraints:

- Relational Model has 3 basic constraint types that are supported in SQL:
  - **Key constraint:** A primary key value cannot be duplicated
  - **Entity Integrity Constraint:** A primary key value cannot be null
  - **Referential integrity constraints :** The “foreign key” must have a value that is already present as a primary key, or may be null.



# Specifying Attribute Constraints

Other Restrictions on attribute domains:

- Default value of an attribute
  - **DEFAULT** <value>
  - NULL is not permitted for a particular attribute (**NOT NULL**)
- **CHECK** clause (restrictions on attribute domains or individual tuples)
  - CREATE DOMAIN D\_NUM AS INTEGER **CHECK** (D\_NUM > 0 AND D\_NUM < 21);
  - Dnumber INT NOT NULL **CHECK** (Dnumber > 0 AND Dnumber < 21);

# Specifying Key and Referential Integrity Constraints

## ■ PRIMARY KEY clause

- Specifies one or more attributes that make up the primary key of a relation
- Dnumber INT **PRIMARY KEY**;

## ■ UNIQUE clause

- Specifies alternate (secondary) keys (called CANDIDATE keys in the relational model).
- Dname VARCHAR(15) **UNIQUE**;

# Specifying Key and Referential Integrity Constraints (cont'd.)

- **FOREIGN KEY** clause
  - Default operation: reject update on violation
  - Attach **referential triggered action** clause
    - Options include **SET NULL**, **CASCADE**, and **SET DEFAULT**
    - Action taken by the DBMS for **SET NULL** or **SET DEFAULT** is the same for both **ON DELETE** and **ON UPDATE**
    - **CASCADE** option suitable for “relationship” relations, for relations that represent **multivalued attributes**, and for relations that represent **weak entity types**

# Giving Names to Constraints

- Using the Keyword **CONSTRAINT**
  - Name a constraint
  - Useful for later altering
    - Use to identify a particular constraint to be dropped later and replaced with another constraint

# Default attribute values and referential integrity triggered action specification (Fig. 6.2)

```
CREATE TABLE EMPLOYEE
(
    ...,
    Dno          INT          NOT NULL      DEFAULT 1,
    CONSTRAINT EMPPK
        PRIMARY KEY (Ssn),
    CONSTRAINT EMPSUPERFK
        FOREIGN KEY (Super_ssn) REFERENCES EMPLOYEE(Ssn)
            ON DELETE SET NULL      ON UPDATE CASCADE,
    CONSTRAINT EMPDEPTFK
        FOREIGN KEY (Dno) REFERENCES DEPARTMENT(Dnumber)
            ON DELETE SET DEFAULT   ON UPDATE CASCADE);

CREATE TABLE DEPARTMENT
(
    ...,
    Mgr_ssn CHAR(9)          NOT NULL      DEFAULT '888665555',
    ...,
    CONSTRAINT DEPTPK
        PRIMARY KEY (Dnumber),
    CONSTRAINT DEPTSK
        UNIQUE (Dname),
    CONSTRAINT DEPTMGRFK
        FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn)
            ON DELETE SET DEFAULT   ON UPDATE CASCADE);

CREATE TABLE DEPT_LOCATIONS
(
    ...,
    PRIMARY KEY (Dnumber, Dlocation),
    FOREIGN KEY (Dnumber) REFERENCES DEPARTMENT(Dnumber)
        ON DELETE CASCADE         ON UPDATE CASCADE);
```

# Specifying Constraints on Tuples Using CHECK

- Additional Constraints on individual **tuples** within a relation are also possible using CHECK
- CHECK clauses at the end of a CREATE TABLE statement
  - Apply to each tuple individually and check whenever a tuple is inserted or modified
  - Example: at the end of the CREATE TABLE statement for the DEPARTMENT table to make sure ...
    - `CHECK (Dept_create_date <= Mgr_start_date);`
  - CHECK can be used to specify more general constraints using the **CREATE ASSERTION** statement of SQL

# Examples of CHECK Constraints

```
CREATE TABLE Persons (  
    ID int NOT NULL,  
    LastName varchar(255) NOT NULL,  
    FirstName varchar(255),  
    Age int,  
    CHECK (Age>=18)  
);
```

```
CREATE TABLE Persons (  
    ID int NOT NULL,  
    LastName varchar(255) NOT NULL,  
    FirstName varchar(255),  
    Age int CHECK (Age>=18)  
);
```

```
CREATE TABLE Persons (  
    ID int NOT NULL,  
    LastName varchar(255) NOT NULL,  
    FirstName varchar(255),  
    Age int,  
    City varchar(255),  
    CONSTRAINT CHK_Person CHECK (Age>=18 AND City='Sandnes')  
);
```

# Basic Retrieval Queries in SQL

- **SELECT statement**
  - One basic statement for retrieving information from a database
- SQL allows a table to have two or more tuples that are **identical** in all their attribute values
  - Unlike relational model (relational model is strictly set-theory based)
  - Multiset or bag behavior
  - **Tuple-id** may be used as a **key**
  - SQL relations can be constrained to be **sets** by specifying PRIMARY KEY or UNIQUE attributes, or by using the **DISTINCT** option in a query



# Retrieval Queries in SQL (contd.)

- A **bag** or **multi-set** is like a set, but an element may appear more than once.
  - Example:  $\{A, B, C, A\}$  is a bag.  $\{A, B, C\}$  is also a bag that also is a set.
  - Bags also resemble lists, but the order is irrelevant in a bag.
- Example:
  - $\{A, B, A\} = \{B, A, A\}$  as bags
  - However,  $[A, B, A]$  is not equal to  $[B, A, A]$  as lists

# The SELECT-FROM-WHERE Structure of Basic SQL Queries

## ■ Basic form of the `SELECT` statement:

```
SELECT    <attribute list>  
FROM      <table list>  
WHERE     <condition>;
```

where

- <attribute list> is a list of attribute names whose values are to be retrieved by the query.
- <table list> is a list of the relation names required to process the query.
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.

# The SELECT-FROM-WHERE Structure of Basic SQL Queries (cont'd.)

- Logical comparison operators
  - =, <, <=, >, >=, and <>
- **Projection attributes**
  - Attributes whose values are to be retrieved
- **Selection condition**
  - Boolean condition that must be true for any retrieved tuple. Selection conditions include join conditions (see Ch.8) when multiple relations are involved.

# Basic Retrieval Queries

<u>Bdate</u>	<u>Address</u>
1965-01-09	731 Fondren, Houston, TX

<u>Fname</u>	<u>Lname</u>	<u>Address</u>
John	Smith	731 Fondren, Houston, TX
Franklin	Wong	638 Voss, Houston, TX
Ramesh	Narayan	975 Fire Oak, Humble, TX
Joyce	English	5631 Rice, Houston, TX

**Query 0.** Retrieve the birth date and address of the employee(s) whose name is 'John B. Smith'.

Q0:     **SELECT**     Bdate, Address  
         **FROM**     EMPLOYEE  
         **WHERE**     Fname='John' AND Minit='B' AND Lname='Smith';

selection condition

**Query 1.** Retrieve the name and address of all employees who work for the 'Research' department.

Q1:     **SELECT**     Fname, Lname, Address  
         **FROM**     EMPLOYEE, DEPARTMENT  
         **WHERE**     Dname='Research' AND Dnumber=Dno;

join condition

selection  
condition

(Dname='Research') is a **selection condition** (corresponds to a **SELECT** operation in relational algebra)  
(Dnumber=Dno) is a **join condition** (corresponds to a **JOIN** operation in relational algebra)  
Such queries are often called **select-project-join** queries

# Basic Retrieval Queries (Contd.)

(c)

<u>Pnumber</u>	<u>Dnum</u>	<u>Lname</u>	<u>Address</u>	<u>Bdate</u>
10	4	Wallace	291Berry, Bellaire, TX	1941-06-20
30	4	Wallace	291Berry, Bellaire, TX	1941-06-20

**Query 2.** For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birth date.

**Q2:**      **SELECT**      Pnumber, Dnum, Lname, Address, Bdate  
             **FROM**        PROJECT, DEPARTMENT, EMPLOYEE  
             **WHERE**      Dnum=Dnumber **AND** Mgr\_ssn=Ssn **AND**  
                         Plocation='Stafford';

- In Q2, there are **two join conditions**
- The **join condition** Dnum=Dnumber relates a project to its controlling department
- The **join condition** Mgr\_ssn=Ssn relates the controlling department to the employee who manages that department

# Ambiguous Attribute Names

- Same name can be used for two (or more) attributes in different relations
  - As long as the attributes are in different relations
  - Must qualify the attribute name with the relation name to prevent ambiguity

```
Q1A:  SELECT  Fname, EMPLOYEE.Name, Address
        FROM    EMPLOYEE, DEPARTMENT
        WHERE   DEPARTMENT.Name='Research' AND
                DEPARTMENT.Dnumber=EMPLOYEE.Dnumber;
```

# Aliasing, and Renaming

- **Aliases or tuple variables**

- Declare **alternative relation names E and S** to refer to the EMPLOYEE relation twice in a query:

- `EMPLOYEE AS E, EMPLOYEE AS S`

- **Query 8.** For each employee, retrieve the employee's first and last name and the first and last name of his or her immediate supervisor.

- `SELECT E.Fname, E.Lname, S.Fname, S.Lname  
FROM EMPLOYEE AS E, EMPLOYEE AS S  
WHERE E.Super_ssn=S.Ssn;`

- Recommended practice to **abbreviate names** and to prefix same or similar attribute from multiple tables.

# Aliasing, Renaming and Tuple Variables (contd.)

- The attribute names can also be renamed
  - `EMPLOYEE AS E(Fn, Mi, Ln, Ssn, Bd, Addr, Sex, Sal, Sssn, Dno)`
- Note that the relation EMPLOYEE now has a variable name E which corresponds to a **tuple variable**
  - The alternate relation names E and S are called aliases or tuple variables for the EMPLOYEE relation
  - We can think of E and S as **two different copies** of EMPLOYEE; E represents employees in role of *supervisees* and S represents employees in role of *supervisors*
- The keyword “AS” may be dropped in most SQL implementations (e.g., EMPLOYEE E, EMPLOYEE S)



# Unspecified WHERE Clause and Use of the Asterisk

- Missing WHERE clause
  - Indicates **no condition** on tuple selection
  - This is equivalent to the condition **WHERE TRUE**
- Effect is a **CROSS PRODUCT**
  - **no join condition**
  - Result is **all possible tuple combinations** (or the Algebra operation of Cartesian Product– see Ch.8)

Queries 9 and 10. Select all EMPLOYEE Ssns (Q9) and all combinations of EMPLOYEE Ssn and DEPARTMENT Dname (Q10) in the database.

Q9:     **SELECT**     Ssn  
          **FROM**     EMPLOYEE;

Q10:    **SELECT**     Ssn, Dname  
          **FROM**     EMPLOYEE, DEPARTMENT;

# Unspecified WHERE Clause and Use of the Asterisk (cont'd.)

- Specify an asterisk (\*)
  - Retrieve **all the attribute values** of the selected tuples
  - The \* can be prefixed by the relation name; e.g., **EMPLOYEE.\*** refers to all attributes of the EMPLOYEE table

Q1C:    SELECT    \*  
         FROM     EMPLOYEE  
         WHERE    Dno=5;

Q1D:    SELECT    \*  
         FROM     EMPLOYEE, DEPARTMENT  
         WHERE    Dname='Research' AND Dno=Dnumber;

Q10A:   SELECT    \*  
         FROM     EMPLOYEE, DEPARTMENT;

# Tables as Sets in SQL

- SQL does **not** automatically **eliminate duplicate tuples** in query results
  - Duplicate elimination is an **expensive** operation (sort and eliminate)
  - User may want to see duplicate tuples in the query results
  - For **aggregate operations** (See sec 7.1.7) duplicates must be accounted for
- Use the keyword **DISTINCT** in the **SELECT** clause
  - **Only distinct tuples should remain in the result**

Query 11. Retrieve the salary of every employee (Q11) and all distinct salary values (Q11A).

Q11:     **SELECT**     **ALL** Salary  
         **FROM**     **EMPLOYEE**;

Q11A:   **SELECT**   **DISTINCT** Salary  
         **FROM**   **EMPLOYEE**;

# Tables as Sets in SQL (cont'd.)

## ■ Set operations

- **UNION, EXCEPT** (difference), **INTERSECT**
- Corresponding **multiset** operations: UNION ALL, EXCEPT ALL, INTERSECT ALL)
- Type compatibility is needed for these operations to be valid

**Query 4.** Make a list of all project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.

```
Q4A: (SELECT DISTINCT Pnumber
      FROM PROJECT, DEPARTMENT, EMPLOYEE
      WHERE Dnum=Dnumber AND Mgr_ssn=Ssn
            AND Lname='Smith' )

      UNION
      ( SELECT DISTINCT Pnumber
        FROM PROJECT, WORKS_ON, EMPLOYEE
        WHERE Pnumber=Pno AND Essn=Ssn
              AND Lname='Smith' );
```

# Set Operations

- The resulting relations of these set operations are **sets** of tuples
  - duplicate tuples are eliminated from the result
- The set operations apply only to union compatible relations
  - The two relations must have the same attributes and the attributes must appear in the same order
- SQL also has corresponding multiset operations using keyword **ALL**
  - UNION ALL, EXCEPT ALL, INTERSECT ALL

# Substring Pattern Matching and Arithmetic Operators

- **LIKE** comparison operator
  - Used for string **pattern matching**
  - **%** replaces an arbitrary number of **zero or more characters**
  - underscore (**\_**) replaces a **single character**
  - Examples:
    - **WHERE Address LIKE '%Houston,TX%';**
    - **WHERE Ssn LIKE '\_\_ 1\_\_ 8901';**
- If **'\_'** or **'%'** is needed as a literal character in the string, the character should be preceded by an escape character, which is specified after the string using the keyword ESCAPE
  - **'AB\\_CD\%EF' ESCAPE '\'** represents the string **'AB\_CD%EF'** because **\** is specified as the **escape character**

# Substring Pattern Matching and Arithmetic Operators (cont'd.)

- Query 12: Retrieve all employees whose address is in Houston, Texas.
  - Here, the value of the ADDRESS attribute must contain the substring 'Houston,TX' in it.

```
Q12:      SELECT      Fname, Lname
           FROM        EMPLOYEE
           WHERE        Address LIKE '%Houston,TX%';
```

# Substring Pattern Matching and Arithmetic Operators (cont'd.)

- Query 12A: Retrieve all employees who were born during the 1950s.
  - Here, '5' must be the third character of the string (according to our format for date), so the BDATE value is ' \_ 5 \_ \_ \_ \_ \_ ', with each underscore as a place holder for a single arbitrary character.

```
Q12A:      SELECT      Fname, Lname
            FROM        EMPLOYEE
            WHERE        Bdate LIKE ' _ 5 _ _ _ _ _ ';
```



# Arithmetic Operations

- Standard arithmetic operators:
  - Addition (+), subtraction (−), multiplication (\*), and division (/) may be **included as a part of SELECT**
- **Query 13.** Show the resulting salaries if every employee working on the 'ProductX' project is given a 10 percent raise.

```
SELECT E.Fname, E.Lname, 1.1 * E.Salary AS Increased_sal  
FROM EMPLOYEE AS E, WORKS_ON AS W, PROJECT AS P  
WHERE E.Ssn=W.Essn AND W.Pno=P.Pnumber AND  
P.Pname='ProductX';
```

# Substring Pattern Matching and Arithmetic Operators (cont'd.)

- **BETWEEN** comparison operator
- Query 14: Retrieve all employee in department 5 whose salary is between \$30,000 and \$40,000

```
Q14:      SELECT      *  
          FROM        EMPLOYEE  
          WHERE        (Salary BETWEEN 30000 AND 40000)  
                     AND Dno=5;
```

- The condition (Salary BETWEEN 30000 AND 40000) is equivalent to ((Salary>=30000) AND (Salary<=40000))

# Ordering of Query Results

- Use **ORDER BY** clause
  - Used to sort the tuples in a query result based on the values of some attribute(s)
  - The default order is in **ascending order** of values
  - Keyword **DESC** to see result in a descending order of values
  - Keyword **ASC** to specify ascending order explicitly
  - Typically placed at the end of the query
  - Ex: if we want descending order on Dname and ascending order on Lname, Fname
    - ORDER BY D.Dname DESC, E.Lname ASC, E.Fname ASC

# Ordering of Query Results (cont'd.)

- Query 15: Retrieve a list of employees and the projects each works in, ordered by the employee's department, and within each department ordered alphabetically by employee last name, first name.

```
Q15:      SELECT      Dname, Lname, Fname, Pname
           FROM        DEPARTMENT, EMPLOYEE,
           WORKS_ON, PROJECT
           WHERE        Dnumber=Dno AND Ssn=Essn
           AND Pno=Pnumber
           ORDER BY     Dname, Lname, Fname;
```

# Basic SQL Retrieval Query Block

```
SELECT    <attribute list>
FROM      <table list>
[ WHERE   <condition> ]
[ ORDER BY <attribute list> ];
```

- The SELECT-clause lists the **attributes** or **functions** to be retrieved
- The FROM-clause specifies all **relations** (or **aliases**) needed in the simple query
- The WHERE-clause specifies the **conditions for selection** and **join** of tuples from the relations specified in the FROM-clause
- ORDER BY specifies an **order** for displaying the result of a query

# INSERT, DELETE, and UPDATE Statements in SQL

- Three commands used to modify the database:
  - INSERT, DELETE, and UPDATE
- INSERT typically inserts a tuple (row) in a relation (table)
- UPDATE may update a number of tuples (rows) in a relation (table) that **satisfy the condition**
- DELETE may also update a number of tuples (rows) in a relation (table) that **satisfy the condition**

# INSERT

- In its simplest form, it is used to add one or more tuples to a relation
- Attribute values should be listed in the same order as the **attributes** were specified in the **CREATE TABLE** command
- Constraints on data types are observed automatically
- **Any integrity constraints** as a part of the DDL specification **are enforced**

# The INSERT Command

- Specify the relation name and a list of values for the tuple. **All values including nulls are supplied.**

```
U1:    INSERT INTO    EMPLOYEE
        VALUES      ( 'Richard', 'K', 'Marini', '653298653', '1962-12-30', '98
                        Oak Forest, Katy, TX', 'M', 37000, '653298653', 4 );
```

- The variation below inserts multiple tuples where a new table is loaded **values from the result of a query.**

```
U3B:    INSERT INTO    WORKS_ON_INFO ( Emp_name, Proj_name,
        Hours_per_week )
        SELECT          E.Lname, P.Pname, W.Hours
        FROM             PROJECT P, WORKS_ON W, EMPLOYEE E
        WHERE            P.Pnumber=W.Pno AND W.Essn=E.Ssn;
```



# INSERT (cont'd.)

- An alternate form of INSERT specifies explicitly the attribute names that correspond to the values in the new tuple
  - Attributes with NULL values can be left out
- Example: Insert a tuple for a new EMPLOYEE for whom we only know the Fname, Lname, Dno, and Ssn attributes.

```
U1A: INSERT INTO EMPLOYEE (Fname, Lname, Dno, Ssn)
      VALUES ('Richard', 'Marini', 4, '653298653');
```

# INSERT (cont'd.)

- Important Note: Only the constraints specified in the DDL commands are automatically enforced by the DBMS when updates are applied to the database
  - It is the responsibility of user to check that any such constraints whose checks are not implemented by the DBMS are not violated

U2: INSERT INTO EMPLOYEE (Fname, Lname, Ssn, Dno)  
VALUES ('Robert', 'Hatcher', '980760540', 2);

- U2 is rejected if referential integrity checking is provided by DBMS (assume that no DEPARTMENT tuple with Dnumber=2 exists)

U2A: INSERT INTO EMPLOYEE (Fname, Lname, Dno)  
VALUES ('Robert', 'Hatcher', 5);

- U2A is rejected if NOT NULL checking is provided by DBMS (assume that there is a NOT NULL constraint on Ssn)

# BULK LOADING OF TABLES

- Another variation of **INSERT** is used for **bulk-loading** of several tuples into tables
- A new table TNEW can be created with the same attributes as T and load some of the data currently in T into TNEW using LIKE and DATA in the syntax, it can be loaded with entire data.
- **EXAMPLE:**

```
CREATE TABLE D5EMPS LIKE EMPLOYEE  
(SELECT E.*  
FROM EMPLOYEE AS E  
WHERE E.Dno=5) WITH DATA;
```

# DELETE

- Removes tuples from a relation
  - Includes a WHERE-clause to select the tuples to be deleted
  - Referential integrity should be enforced
  - Tuples are deleted from *only one table at a time* (unless **CASCADE** is specified on a referential integrity constraint)
    - The deletion may propagate to tuples in other relations if *referential triggered actions* are specified in the referential integrity constraints of the DDL
  - A **missing WHERE-clause** specifies that *all tuples* in the relation are **to be deleted**; the table then becomes an empty table
  - The number of tuples deleted depends on the number of tuples in the relation that **satisfy the WHERE-clause**

# The DELETE Command

- Removes tuples from a relation
  - Includes a `WHERE` clause to select the tuples to be deleted. The number of tuples deleted will vary.

U4A:	DELETE FROM	EMPLOYEE
	WHERE	Lname='Brown';
U4B:	DELETE FROM	EMPLOYEE
	WHERE	Ssn='123456789';
U4C:	DELETE FROM	EMPLOYEE
	WHERE	Dno=5;
U4D:	DELETE FROM	EMPLOYEE;

# UPDATE

- Used to modify attribute values of one or more selected tuples
- A **WHERE-clause** selects the tuples to be modified
- An additional **SET-clause** specifies the attributes to be modified and their new values
- Each command modifies tuples *in the same relation*
- **Referential integrity** specified as part of DDL specification is **enforced**

# UPDATE (contd.)

- Example: Change the location and controlling department number of project number 10 to 'Bellaire' and 5, respectively

```
U5:      UPDATE      PROJECT
          SET         PLOCATION = 'Bellaire',
                   DNUM = 5
          WHERE       PNUMBER=10
```

# UPDATE (contd.)

- Example: Give all employees in the 'Research' department a 10% raise in salary.

```
U6:UPDATE      EMPLOYEE
      SET       SALARY = SALARY *1.1
      WHERE     DNO IN (SELECT DNUMBER
                          FROM   DEPARTMENT
                          WHERE   DNAME='Research')
```

- In this request, the modified SALARY value depends on the original SALARY value in each tuple
  - The reference to the SALARY attribute on the right of = refers to the old SALARY value **before** modification
  - The reference to the SALARY attribute on the left of = refers to the new SALARY value **after** modification



# Additional Features of SQL

- Techniques for specifying complex retrieval queries (see Ch.7)
- Writing programs in various programming languages that include SQL statements: Embedded and dynamic SQL, SQL/CLI (Call Level Interface) and its predecessor ODBC, SQL/PSM (Persistent Stored Module) (See Ch.10)
- Set of commands for specifying physical database design parameters, file structures for relations, and access paths, e.g., CREATE INDEX

# Additional Features of SQL (cont'd.)

- Transaction control commands (Ch.20)
- Specifying the granting and revoking of privileges to users (Ch.30)
- Constructs for creating triggers (Ch.26)
- Enhanced relational systems known as object-relational define relations as classes. Abstract data types (called User Defined Types- UDTs) are supported with CREATE TYPE
- New technologies such as XML (Ch.13) and OLAP (Ch.29) are added to versions of SQL

# Summary

- SQL
  - A Comprehensive language for relational database management
  - Data definition, queries, updates, constraint specification, and view definition
- Covered :
  - Data definition commands for creating tables
  - Commands for constraint specification
  - Simple retrieval queries
  - Database update commands